





```
In [81]: >>> cv_tfidf_train_test(dataframe,label,vectorizer,ngram):
# Split the data into X and y data sets
X = dataframe.comment_text
y = dataframe[label]

# Split our data into training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=50)

# Using vectorizer and removing stopwords
cv1 = vectorizer(ngram_range=(ngram), stop_words='english')

# Transforming x_train and x_test
X_train_cv1 = cv1.fit_transform(X_train)
X_test_cv1 = cv1.transform(X_test)

## Machine learning models

## Logistic regression
lr = LogisticRegression()
lr.fit(X_train_cv1, y_train)

## k-nearest neighbours
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_cv1, y_train)

## Naive Bayes
bnb = BernoulliNB()
bnb.fit(X_train_cv1, y_train)

## Multinomial naive bayes
mnb = MultinomialNB()
mnb.fit(X_train_cv1, y_train)

## Support vector machine
svm_model = LinearSVC()
svm_model.fit(X_train_cv1, y_train)

## Random Forest
randomforest = RandomForestClassifier(n_estimators=100, random_state=50)
randomforest.fit(X_train_cv1, y_train)

f1_score_data = {'F1 Score':(f1_score(lr.predict(X_test_cv1), y_test), f1_score(knn.predict(X_test_cv1), y_test), f1_score(bnb.predict(X_test_cv1), y_test), f1_score(mnb.predict(X_test_cv1), y_test), f1_score(svm_model.predict(X_test_cv1), y_test), f1_score(randomforest.predict(X_test_cv1), y_test))
# Saving f1 score results into a dataframe
df_f1 = pd.DataFrame(f1_score_data, index=['Log Regression','KNN', 'BernoulliNB', 'MultinomialNB', 'SVM', 'Random Forest'])

return df_f1

In [ ]:

In [ ]:

In [ ]:

In [ ]:
```

Evaluating model performance using evaluation metrics.

```
In [82]: severe_toxic_comment_cv = cv_tfidf_train_test(severe_toxic_comment_balanced, 'severe_toxic', TfidfVectorizer, (1,1))
severe_toxic_comment_cv.rename(columns={'F1 Score': 'F1 Score(severe_toxic)'}, inplace=True)
severe_toxic_comment_cv
# Multinomial NB has higher F1 score
```

Out [82]:

F1 Score(severe_toxic)	
Log Regression	0.940282
KNN	0.860192
BernoulliNB	0.790738
MultinomialNB	0.932377
SVM	0.937901
Random Forest	0.941176

In [ ]:

```
In [83]: obscene_comment_cv = cv_tfidf_train_test(obscene_comment_balanced, 'obscene', TfidfVectorizer, (1,1))
obscene_comment_cv.rename(columns={'F1 Score': 'F1 Score(obscene)'}, inplace=True)
obscene_comment_cv
# Random Forest has higher F1 score
```

Out [83]:

F1 Score(obscene)	
Log Regression	0.901183
KNN	0.625341
BernoulliNB	0.766640
MultinomialNB	0.887496
SVM	0.915613
Random Forest	0.884261

```
In [84]: threat_comment_cv = cv_tfidf_train_test(threatening_comment_balanced, 'threat', TfidfVectorizer, (1,1))
threat_comment_cv.rename(columns={'F1 Score': 'F1 Score(threat)'}, inplace=True)
threat_comment_cv
# Random Forest has higher F1 score
```

Out [84]:

F1 Score(threat)	
Log Regression	0.897338
KNN	0.852459
BernoulliNB	0.745205
MultinomialNB	0.902098
SVM	0.894737
Random Forest	0.923077

```
In [85]: insult_comment_cv = cv_tfidf_train_test(insulting_comment_balanced, 'insult', TfidfVectorizer, (1,1))
insult_comment_cv.rename(columns={'F1 Score': 'F1 Score(insult)'}, inplace=True)
insult_comment_cv
# SVM has higher F1 score
```

Out [85]:

F1 Score(insult)	
Log Regression	0.901851
KNN	0.320661
BernoulliNB	0.776986
MultinomialNB	0.896299
SVM	0.906218
Random Forest	0.890821

```
In [86]: identity_hatecomment_cv = cv_tfidf_train_test(identity_hate_comment_balanced, 'identity_hate', TfidfVectorizer, (1,1))
identity_hatecomment_cv.rename(columns={'F1 Score': 'F1 Score(identity_hate)'}, inplace=True)
identity_hatecomment_cv
# MultinomialNB has higher F1 score
```

Out [86]:

F1 Score(identity_hate)	
Log Regression	0.905707
KNN	0.820046
BernoulliNB	0.776699
MultinomialNB	0.903302
SVM	0.898066
Random Forest	0.888087

```
In [87]: X = Toxic comment balanced comment text
y = Toxic comment balanced['toxic']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Initiate a Tfidf vectorizer
tfv = TfidfVectorizer(ngram_range=(1,1), stop_words='english')

X_train_fit = tfv.fit_transform(X_train)
X_test_fit = tfv.transform(X_test)
randomforest = RandomForestClassifier(n_estimators=100, random_state=50)
randomforest.fit(X_train_fit, y_train)
randomforest.predict(X_test_fit)
```

Out [87]: array([0, 1, 1, ..., 1, 1, 1], dtype=int64)

```
In [88]: ## Testing the model to check if the given text is toxic or not.
```

```
In [89]: comment1 = ['i killed an insect and ate it']
comment1_vect = tfv.transform(comment1)
randomforest.predict_proba(comment1_vect)[:,1]
## As seen below the above comment is 73 percent toxic
```

Out [89]: array([0.73519444])

```
In [90]: comment2 = ['is this sentence a good one']
comment2_vect = tfv.transform(comment2)
randomforest.predict_proba(comment2_vect)[:,1]
## As seen below the above comment is 0.08 percent toxic which says the comment is not toxic
```

Out [90]: array([0.08770635])

```
In [91]: comment2 = ['truth will prevail']
comment2_vect = tfv.transform(comment2)
randomforest.predict_proba(comment2_vect)[:,1]
## The above comment is 46 percent toxic.
```

Out [91]: array([0.46238997])

```
In [ ]:
```